

Homework 4

1

```
library('caret')
return_probs <- function(row){
  row <- as.numeric(row)

  product_1 <- p_v1_v11[,1][row[1]+1]* p_v2_v11[,1][row[2]+1]* p_v3_v11[,1][row[3]+1]* p_v4_v11[,1][row[4]+1]
  product_2 <- p_v1_v11[,2][row[1]+1]* p_v2_v11[,2][row[2]+1]* p_v3_v11[,2][row[3]+1]* p_v4_v11[,2][row[4]+1]
  product_3 <- p_v1_v11[,3][row[1]+1]* p_v2_v11[,3][row[2]+1]* p_v3_v11[,3][row[3]+1]* p_v4_v11[,3][row[4]+1]

  c(P_1*product_1, P_2*product_2, P_3*product_3)
}

binary_data <- read.csv('naive_bayes_binary.csv')
binary_data_train <- binary_data[1:(length(binary_data$V1)/2),]
binary_data_test <- binary_data[2501:5000,]
true_values <- binary_data_test$V1
binary_data_test <- binary_data_test[1:10]

length(binary_data_train$V1)

## [1] 2500

#create tables of all variables and derive their conditional probs

table_v1_v11 <- table(binary_data_train$V1, binary_data_train$V11)
table_v2_v11 <- table(binary_data_train$V2, binary_data_train$V11)
table_v3_v11 <- table(binary_data_train$V3, binary_data_train$V11)
table_v4_v11 <- table(binary_data_train$V4, binary_data_train$V11)
table_v5_v11 <- table(binary_data_train$V5, binary_data_train$V11)
table_v6_v11 <- table(binary_data_train$V6, binary_data_train$V11)
table_v7_v11 <- table(binary_data_train$V7, binary_data_train$V11)
table_v8_v11 <- table(binary_data_train$V8, binary_data_train$V11)
table_v9_v11 <- table(binary_data_train$V9, binary_data_train$V11)
table_v10_v11 <- table(binary_data_train$V10, binary_data_train$V11)
```

```

# table_vecs <-

p_v1_v11 <- prop.table(table_v1_v11,2)
p_v2_v11 <- prop.table(table_v2_v11,2)
p_v3_v11 <- prop.table(table_v3_v11,2)
p_v4_v11 <- prop.table(table_v4_v11,2)
p_v5_v11 <- prop.table(table_v5_v11,2)
p_v6_v11 <- prop.table(table_v6_v11,2)
p_v7_v11 <- prop.table(table_v7_v11,2)
p_v8_v11 <- prop.table(table_v8_v11,2)
p_v9_v11 <- prop.table(table_v9_v11,2)
p_v10_v11 <- prop.table(table_v10_v11,2)

number_1 <- sum(binary_data_train$V11 ==1)
number_2 <- sum(binary_data_train$V11 ==2)
number_3 <- sum(binary_data_train$V11 ==3)

P_1 <- number_1/length(binary_data_train$V11)
P_2 <- number_2/length(binary_data_train$V11)
P_3 <- number_3/length(binary_data_train$V11)

probs_v11_numbers <- c(P_1, P_2, P_3)

results <- c()
for (i in 1:length(binary_data_test$V1)) {
  temp <- return_probs(binary_data_test[i,])
  results <- c(results, match(max(temp), temp))
}

# preds <- return_probs(binary_data[1,])

mean(results == true_values)

## [1] 0.9424

confusion <- table(results, true_values)
confusion

##           true_values
## results      1      2      3
##      1  177    16     6
##      2   23   206    13
##      3    55    31  1973

```

2

a)

```
library('rpart')
library('rpart.plot')
student_data <- read.csv('student/student-mat.csv', sep=';')

student_data$class <- ifelse(student_data$G3 > 10,0,1)
new_dat <- student_data[1:30]
new_dat <- cbind(new_dat, student_data[34])

sample_ind <- sample(nrow(new_dat),nrow(new_dat)*0.70)
train <- new_dat[sample_ind,]
test <- new_dat[-sample_ind,]

set.seed(100)
fit <- rpart(formula = class~., data = train, method = "class", control = rpart.control(cp = 0.01))

printcp(fit)

##
## Classification tree:
## rpart(formula = class ~ ., data = train, method = "class", control = rpart.control(cp = 0.01))
##
## Variables actually used in tree construction:
## [1] age      Dalc      failures famrel  famsize  Fedu      Fjob
## [8] freetime goout    guardian health  Mjob      reason   schoolsup
## [15] Walc
##
## Root node error: 134/276 = 0.48551
##
## n= 276
##
##          CP nsplit rel error  xerror    xstd
## 1  0.2537313      0  1.00000  1.08209  0.061910
## 2  0.0373134      1  0.74627  0.84328  0.060964
## 3  0.0223881      6  0.53731  0.94030  0.061755
## 4  0.0186567      7  0.51493  0.90299  0.061518
## 5  0.0111940      9  0.47761  0.88806  0.061399
## 6  0.0099502     13  0.41791  0.89552  0.061460
## 7  0.0074627     16  0.38806  0.89552  0.061460
## 8  0.0000000     18  0.37313  0.91791  0.061622
## 9 -1.0000000     23  0.37313  0.91791  0.061622

test$pred <- predict(fit, test, type = "class")
```

```

#pruning
fit_pruned <- prune(fit, cp = 0.038)
test$pred <- predict(fit_pruned, test, type = "class")

rpart.plot(fit_pruned)

```




Figure 1: plot of chunk unnamed-chunk-2

b)

training error = root node error X relative error

training error = $0.47 * 0.59$

generalisation error = $0.47 * 0.84$

c)

```
imortant <- fit_pruned$variable.importance
```

```
imortant
```

```
##   failures      age  guardian      higher  absences
## 13.8954468  2.3957667  0.7187300  0.7187300  0.4791533
```

According to the given output 'failures' is the most important variable.

d)

```
student_data_reg <- read.csv('student/student-mat.csv', sep=';')
```

```
new_dat <- student_data[1:30]
```

```
new_dat <- cbind(new_dat, student_data[33])
```

```
sample_ind <- sample(nrow(new_dat), nrow(new_dat)*0.70)
```

```
train_reg <- new_dat[sample_ind,]
```

```
test_reg <- new_dat[-sample_ind,]
```

```
set.seed(100)
```

```
fit_reg <- rpart(formula = G3~., data = train_reg, method = "anova", control = rpart.control
```

```
printcp(fit_reg)
```

```
##
```

```
## Regression tree:
```

```
## rpart(formula = G3 ~ ., data = train_reg, method = "anova", control = rpart.control(cp =
```

```
##
```

```
## Variables actually used in tree construction:
```

```
## [1] absences activities age failures famrel Fedu
```

```
## [7] Fjob freetime goout Medu Mjob reason
```

```
## [13] sex studytime Walc
```

```
##
```

```
## Root node error: 5280.8/276 = 19.133
```

```
##
```

```
## n= 276
```

```
##
```

```
##          CP nsplit rel error  xerror    xstd
```

```
## 1  0.1235880      0  1.00000 1.00392 0.097113
```

```
## 2  0.0762291      1  0.87641 0.90362 0.090044
```

```
## 3  0.0258000      2  0.80018 0.81242 0.086316
```

```
## 4  0.0234981      4  0.74858 0.91164 0.095049
```


```
## 5  0.0192030      7  0.67809 0.90836 0.096999
```

```
## 6  0.0153929      8  0.65889 0.89376 0.096556
```

```
## 7  0.0115373      9  0.64349 0.90927 0.097631
## 8  0.0112641     10  0.63196 0.90675 0.095669
## 9  0.0101256     11  0.62069 0.90681 0.095764
## 10 0.0094987     12  0.61057 0.90961 0.095318
## 11 0.0088311     13  0.60107 0.92750 0.097520
## 12 0.0087288     14  0.59224 0.93772 0.099166
## 13 0.0084118     15  0.58351 0.93772 0.099166
## 14 0.0074126     16  0.57510 0.95692 0.100647
## 15 0.0072139     18  0.56027 0.95571 0.099703
## 16 0.0068942     19  0.55306 0.95707 0.099689
## 17 0.0056967     20  0.54616 0.95704 0.099669
## 18 0.0055199     21  0.54047 0.95549 0.099844
## 19 -1.0000000     22  0.53495 0.95549 0.099844
```

```
fit_pruned_reg <- prune(fit_reg, cp = 0.027)
```

```
rpart.plot(fit_pruned_reg)
```



training error = root node error X relative error

training error = $19.33 * 0.8 = 15.65$

generalisation error = $19.33 * 0.81 = 15.6573$

```
important2 <- fit_pruned_reg$variable.importance
important2
```

```
##   failures   absences      age   guardian traveltime    higher
## 652.64610 402.55244 176.98877 110.61798   47.35911   44.24719
##      goout
## 23.67956
```

Here too 'failures' is the most important variable.

4

a)

```
strange_data <- read.csv('strange_binary.csv')

set.seed(100)
fit_strange <- rpart(c~., data=strange_data, method='class', control = rpart.control(cp = 0))

printcp(fit_strange)

##
## Classification tree:
## rpart(formula = c ~ ., data = strange_data, method = "class",
##       control = rpart.control(cp = 0))
##
## Variables actually used in tree construction:
## [1] X    X.1 X.2 X.3 X.4 X.6 X.7 X.9
##
## Root node error: 64/200 = 0.32
##
## n= 200
##
##          CP nsplit rel error xerror   xstd
## 1 0.0572917      0  1.00000 1.0000 0.10308
## 2 0.0312500      3  0.82812 1.0938 0.10540
## 3 0.0156250      4  0.79688 1.1094 0.10574
## 4 0.0078125      5  0.78125 1.1094 0.10574
## 5 0.0000000      9  0.75000 1.0781 0.10504

#          CP nsplit rel error xerror   xstd
# 1 0.0572917      0  1.00000 1.0000 0.10308
# 2 0.0312500      3  0.82812 1.0938 0.10540
# 3 0.0156250      4  0.79688 1.1094 0.10574
# 4 0.0078125      5  0.78125 1.1094 0.10574
# 5 0.0000000      9  0.75000 1.0781 0.10504

fit_pruned_strange <- prune(fit_strange, cp = 0.031250)
strange_predict <- predict(fit_pruned_strange, strange_data, type='class')
score <- mean(strange_predict == strange_data$c)
rpart.plot(fit_pruned_strange)

#cp should be 0.00781
#root node error = 0.32
#rel error = 0.828
#xerror = 1.093
```

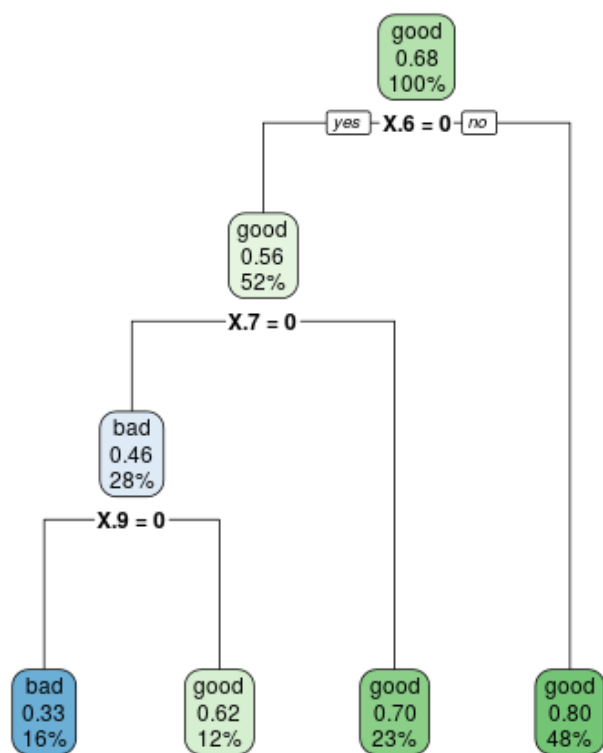



Figure 2: plot of chunk unnamed-chunk-6

Here

Training error = root node error X relative error

Training error = $0.32 * 0.828$

Generalisation error = root node error X xerror

Generalisation error = $0.32 * 1.093$

It is not reasonable to assume error rate on generalised data set will be similar to error rate on test set. Even if the relative error is low, xerror is quite high on this tree. So for test data, the error will be higher.

b)

```
new_feature <- rowSums(strange_data[1:10])
strange_data$new_fet <- new_feature

set.seed(100)
fit_strange_new <- rpart(c~., data=strange_data, method='class', control = rpart.control(cp
printcp(fit_strange_new)

##
## Classification tree:
## rpart(formula = c ~ ., data = strange_data, method = "class",
##       control = rpart.control(cp = 0))
##
## Variables actually used in tree construction:
## [1] new_fet X      X.1      X.5      X.9
##
## Root node error: 64/200 = 0.32
##
## n= 200
##
##          CP nsplit rel error  xerror    xstd
## 1 0.296875     0  1.00000 1.00000 0.103078
## 2 0.009375     1  0.70312 0.70312 0.092274
## 3 0.000000     6  0.65625 0.81250 0.096925

fit_pruned_strange_new <- prune(fit_strange_new, cp = 0.009375)
printcp(fit_pruned_strange_new)

##
## Classification tree:
## rpart(formula = c ~ ., data = strange_data, method = "class",
##       control = rpart.control(cp = 0))
##
## Variables actually used in tree construction:
```

```
## [1] new_fet
##
## Root node error: 64/200 = 0.32
##
## n= 200
##
##          CP nsplit rel error  xerror    xstd
## 1 0.296875      0  1.00000 1.00000 0.103078
## 2 0.009375      1  0.70312 0.70312 0.092274
rpart.plot(fit_pruned_strange_new)
```




Figure 3: plot of chunk unnamed-chunk-7

```
#0.32 = root node error
#rel error = 0.703
#xerror = 0.703
```

Training error = root node error X relative error

Training error = $0.32 * 0.703 = 0.20$

Generalisation error = $0.32 * 0.703$

correct on trainings = $1 - \text{training error} = 1 - 0.22496 = 0.77504$

6

```
classification_data <- read.csv('classification_accuracy.csv')
```

```
wins_decision_tree <- sum(classification_data$decision_tree > classification_data$svm) + sum(classification_data$decision_tree > classification_data$naive_bayes)
wins_svm <- sum(classification_data$svm > classification_data$decision_tree) + sum(classification_data$svm > classification_data$naive_bayes)
wins_naive_bayes <- sum(classification_data$naive_bayes > classification_data$svm) + sum(classification_data$naive_bayes > classification_data$decision_tree)
```

```
loss_decision_tree <- sum(classification_data$decision_tree < classification_data$svm) + sum(classification_data$decision_tree < classification_data$naive_bayes)
loss_svm <- sum(classification_data$svm < classification_data$decision_tree) + sum(classification_data$svm < classification_data$naive_bayes)
loss_naive_bayes <- sum(classification_data$naive_bayes < classification_data$svm) + sum(classification_data$naive_bayes < classification_data$decision_tree)
```

```
draw_decision_tree <- sum(classification_data$decision_tree == classification_data$svm) + sum(classification_data$decision_tree == classification_data$naive_bayes)
draw_svm <- sum(classification_data$svm == classification_data$decision_tree) + sum(classification_data$svm == classification_data$naive_bayes)
draw_naive_bayes <- sum(classification_data$naive_bayes == classification_data$svm) + sum(classification_data$naive_bayes == classification_data$decision_tree)
```

```
a <- c(wins_decision_tree, draw_decision_tree, loss_decision_tree)
b <- c(wins_naive_bayes, draw_naive_bayes, loss_naive_bayes)
c <- c(wins_svm, draw_svm, loss_svm)
```

```
k <- rbind(a,b,c)
rownames(k) = c('Decision_tree', 'Naive_bayes', 'SVM')
colnames(k) = c('Wins', 'Draw', 'Losses')
```

```
confusion_matrix <- k
```

```
confusion_matrix
```

```
##           Wins Draw Losses
## Decision_tree   19    1    26
## Naive_bayes     14    1    31
## SVM             35    0    11
```